
Classification of YouTube Videos using Deep Learning Methods

Jayaram Kuchibhotla
Department of Computer Science
University at Buffalo
Person #:50208766
jayaramk@buffalo.edu

Abstract

This project report describes the methods implemented for YouTube video recognition task. The problem statement is presented by Google Cloud & YouTube-8m video understanding challenge. Apart from the Video Level logistic regression, Frame Level Logistic regression, LSTM (Long Short Term Memory) models provided in the starter code, the newer models such as GRU (Gated Recurrent Unit), BI-LSTM (Bidirectional LSTM), and BI-GRU (Bidirectional GRU) are trained and validated on the Google Cloud Platform. The report has the following steps 1) Introduction 2) Description of the dataset 3) Explanation of above-mentioned methods 4) Discussion of results 5) Problems and Possible Improvements 6) Conclusion

1 Introduction

The deep learning methods are applied successfully in the field of Computer Vision for classifying large-scale image and video datasets such as ImageNet. The objective of automating the task to analyze the content of the videos could possibly be used in applications such as filtering and categorization. As the size of dataset increases, the complexity in recognizing the content of the videos increases. The availability of large video dataset and this public challenge from Google & YouTube helped people to come up with efficient methods in short period. The top scorers in this competition mainly obtained good results using the ensemble of different deep learning methods

2 Description of the dataset

The dataset consists of video data related to areas such as sports, activities, animals, foods, products, tourist attractions, games etc. The aim of the dataset is to understand what is present in each frame of the video just by using the topic label. The number of topics or labels in this dataset is 4716 from 24 diverse categories. The average number of labels per video is 3.4

The dataset is organized as follows:

Partition	Train	Validate	Test	Total
Number of videos	4,906,660	1,401,828	700,640	7,009,128

Each video has both frame level features and video level features.

Frame-level features are stored as [tensorflow.SequenceExample](#) protocol buffers

```
context: {  
  feature: {  
    key : "video_id"  
    value: {  
      bytes_list: {  
        value: [YouTube video id string]  
      }  
    }  
  }  
}  
feature: {  
  key : "labels"
```

```

    value: {
      int64_list: {
        value: [1, 522, 11, 172] # The meaning of the labels can be found here.
      }
    }
  }
}

```

```

feature_lists: {
  feature_list: {
    key : "rgb"
    value: {
      feature: {
        bytes_list: {
          value: [1024 8bit quantized features]
        }
      }
    }
    feature: {
      bytes_list: {
        value: [1024 8bit quantized features]
      }
    }
    ... # Repeated for every second of the video, up to 300
  }
  feature_list: {
    key : "audio"
    value: {
      feature: {
        bytes_list: {
          value: [128 8bit quantized features]
        }
      }
    }
    feature: {
      bytes_list: {
        value: [128 8bit quantized features]
      }
    }
    ... # Repeated for every second of the video, up to 300
  }
}

```

The total size of the frame-level features is 1.71 Terabytes. They are divided into 4096 shards

Video-level features are also stored as [tensorflow.Example](#) protocol buffers

```

features: {
  feature: {
    key : "video_id"
    value: {
      bytes_list: {
        value: [YouTube video id string]
      }
    }
  }
}
feature: {
  key : "labels"
  value: {
    int64_list: {
      value: [1, 522, 11, 172]
    }
  }
}

```

```

    }
  }
}
feature: {
  key : "mean_rgb" # Average of all 'rgb' features for the video
  value: {
    float_list: {
      value: [1024 float features]
    }
  }
}
feature: {
  key : "mean_audio" # Average of all 'audio' features for the video
  value: {
    float_list: {
      value: [128 float features]
    }
  }
}
}
}

```

The total size of the video-level features is 31 Gigabytes. They are divided into 4096 shards

The visual features were extracted by the organizers using [Inception-V3 image annotation model](#), trained on ImageNet. The audio features were extracted using a VGG-inspired acoustic model described in [Hershey et. al.](#) on a preliminary version of YouTube-8M. Both the visual and audio features were PCA-ed and quantized to fit on a single hard disk

3 Explanation of methods

3.1 Video Level Methods

a) Video level Logistic Regression

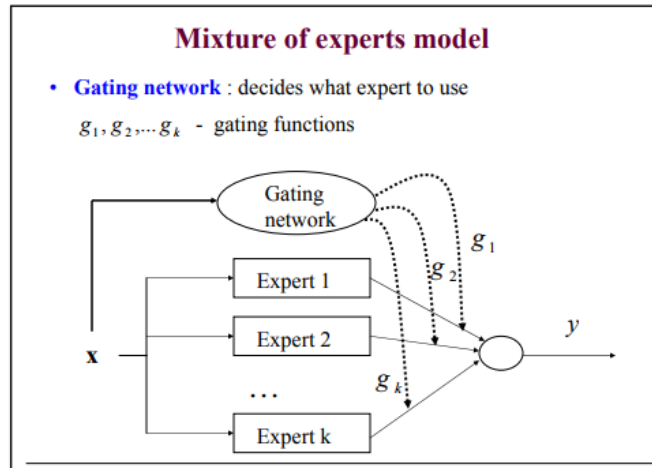
The input video-level features of each video which is a tensor of dimension (batch_size = 128 or 256 , num_features = 1024 or 128 or 1152)) .User can choose to input the video features(length = 1024) or audio features(length = 128) or concatenation of both (length = 1152). Each batch is given as input to the fully connected layer which has sigmoid activation, average pooling, L2 weight regularization with number of classes (vocab size) = 4716 to generate the predictions

3.2 Frame Level methods:

a) Frame level Logistic Regression

For all the frame level models described below the input is a tensor of dimension (batch_size = 128 or 256, num_frames = 300 , num_features = 1024 or 128 or 1152))) .The same strategy as mentioned for the video level logistic regression is used here as well

All the frame level models described below(except the above one) would use 'Mixture of experts' (Moe)(provided in the starter code) in the final/output stage. Mixture of experts is based on the divide-and-conquer principle in which the problem space is divided between a few neural network experts, supervised by a gating network



Learning Moe involves learning the parameters of individual expert networks and parameters of the gating network. Based on the probability values obtained from the gating network the space is divided—partitions belongs to different experts


In this project, the number of mixtures used = 2, gate activations(size = vocab_size * (num_mixtures + 1) and expert activations (size = vocab_size * num_mixtures) are created using the fully connected layers (with l2 regularization) . Gating distribution is formed by using softmax activation of gate activations .Expert distribution is formed by using sigmoid activation of expert activations

Gating and Expert distributions are used to form the final probabilities or prediction values

b) Long short-term memory (LSTM)

LSTM is a recurrent neural network (RNN) architecture that remembers values and used for series classification. Videos consists of variable length images and LSTM is best suited for classifying them

Diagram of a single LSTM cell




In this project (in the starter code),the number of layers are used are two and each layer consists of 1024 cells .

c) Gated Recurrent Unit(GRU)

GRU is also recurrent neural network (RNN) architecture that remembers values and used for series classification. It has fewer parameters compared to LSTM and lacks an output cell. It can be trained relatively fast compared to LSTM network

Gated Recurrent Unit (GRU)

Similar performance as LSTM with less computation.



Cho, Kyunghyun, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. "Learning phrase representations using RNN encoder-decoder for statistical machine translation." AMNLP 2014.

32


In this project (code developed by self),the number of layers are used are two and each layer consists of 1024 cells .

d) Bidirectional Recurrent Neural Networks

- 1) Bidirectional LSTM(BI-LSTM)
- 2) Bidirectional GRU(BI-GRU)

Recurrent neural network (RNNs) also have restrictions, as the future input information cannot be reached from the current state. On the contrary, BRNNs do not require their input data to be fixed. Moreover, their future input information is reachable from the current state. The basic idea of BRNNs is to connect two hidden layers of opposite directions to the same output. By this structure, the output layer can get information from past and future states.

BRNN are especially useful when the context of the input is needed. If the information from either ends of video is used, video context can be easily inferred



In the current project for both BIRNNs (code developed by self), one forward layer and backward layer are used, The states of both forward layer and backward layer are concatenated together to form an output state.

5 Discussion of results

The accuracy of the results can be inferred from the following parameters

1. Hit @1

This is the fraction of test samples that contain at least one of the ground truth labels in the top $k(k=1)$ predictions

2. Precision at Equal Recall Rate

It is video level annotation precision when we retrieve the same number of entities per video as there are in the ground-truth.

3. Global average precision:

The average precision, approximating the area under the precision-recall curve, can then be computed as where P is Precision and R is recall

$$AP = \sum_{j=1}^{10000} P(\tau_j)[R(\tau_j) - R(\tau_{j+1})],$$

The mean/global average precision is computed as the unweighted mean of all the per-class average precisions.

All the performance values range between [0,1] (higher values indicate better performance)

These performance measures are calculated for training and validation datasets only

The results of testing dataset is meant for judgement in the competition .Top 20 label and confidence values are generated in a .csv file for each video or row. It is not feasible to calculate these measures using the starter code for the test dataset.

The snapshots of logs for each one of the 9 variants is as follows

(Owing to the huge size of the dataset,large training duration, Google cloud credit expense) training and validation is restricted to limited number of samples and limited duration

The last line in the log of each snapshot is the final performance measure for that method.

1. Video Level logistic Model (training duration: 42 min 53 sec)

Training:

Cloud ML Job, yt8m_train_20171015_110659	All logs	Any log level	Jump to date
2017-10-15 EDT			
View Options			
11:50:49.581	training step 12169	Loss: 7.05 Examples/sec: 12160.36	
11:50:49.917	training step 12170	Loss: 7.31 Examples/sec: 14353.78 Hit@1: 0.81 PERR: 0.65 GAP: 0.70	

Training accuracy (Hit @1) = 81%

Validation:

Cloud ML Job, yt8m_eval_20171017_133852	All logs	Any log level	Jump to date
2017-10-17 EDT			
View Options			
13:51:41.651	examples_processed: 331776	global_step 12010	Batch Hit@1: 0.800
			Batch PERR: 0.654
			Batch Loss: 7.807
			Examples_per_sec: 3556.187
13:51:42.431	examples_processed: 332800	global_step 12010	Batch Hit@1: 0.791
			Batch PERR: 0.644
			Batch Loss: 7.503
			Examples_per_sec: 3449.787

Validation accuracy (Hit @1) = 79.1%

Testing(predictions.csv ,being shown only for this method) :

Videoid	LabelConfidencePairs
1E+08	1 0.944580 4 0.897149 0 0.699845 2292 0.639824 1833 0.510826 2 0.498553 34 0.433821 297 0.359376 167 0.262460 933 0.241217 198 0.195827 360 0.128280 3547 0.115357 3279 0.093449 19 0.089835 74 0.057152 202
1E+08	77 0.920271 112 0.847571 142 0.819089 21 0.516012 59 0.453498 0 0.111039 8 0.101669 1 0.084844 60 0.066888 57 0.049879 11 0.035834 17 0.034797 4 0.019088 5 0.017757 10 0.015947 262 0.015331 2 0.014108 75 0.0
1E+08	46 0.760294 7 0.562875 3 0.278214 13 0.254399 38 0.128333 130 0.092534 5 0.062423 126 0.030619 150 0.022289 1270 0.013765 1790 0.013600 1 0.012415 4 0.009045 95 0.008160 93 0.006226 16 0.006098 99 0.005230
1E+08	1 0.399384 85 0.162444 62 0.084887 103 0.036210 4 0.031805 27 0.029982 119 0.027678 34 0.024589 64 0.023760 5 0.021960 60 0.021187 8 0.021148 12 0.013430 152 0.009803 40 0.008540 75 0.008249 33 0.008036 47
1.01E+08	7 0.712110 14 0.558702 13 0.244020 24 0.176967 9 0.078686 12 0.076217 1 0.056301 63 0.051904 367 0.035011 3 0.033245 4 0.031652 44 0.029508 39 0.028861 30 0.017998 227 0.016752 11 0.013838 17 0.011994 18 0.0

2. Frame Level Logistic Model with only video features (train duration: 1 hr 34 min)

Training:

Cloud ML Job, yt8m_train_20171015_125643	All logs	Any log level	Jump to date
2017-10-15 EDT			
View Options			
14:26:41.494	training step 16358	Loss: 8.52 Examples/sec: 292.49	
14:26:42.251	training step 16360	Loss: 8.28 Examples/sec: 444.37 Hit@1: 0.80 PERR: 0.61 GAP: 0.67	

Training accuracy (Hit @1) = 80%

Validation:

Cloud ML Job, yt8m_eval_20171017_141220	All logs	Any log level	Jump to date
2017-10-17 EDT			
View Options			
14:42:54.850	examples_processed: 82560	global_step 17291 Batch Hit@1: 0.766 Batch PERR: 0.590 Batch Loss: 8.921 Examples_per_sec: 69.516	
14:42:56.959	examples_processed: 82688	global_step 17291 Batch Hit@1: 0.758 Batch PERR: 0.645 Batch Loss: 7.976 Examples_per_sec: 71.230	

Validation accuracy (Hit @1) = 64.5%

3. Frame Level Logistic Model with video and audio features (4 hr 6 min)

Training:

2017-10-15 EDT			
View Options			
18:08:06.428	training step 38919	Loss: 7.02 Examples/sec: 420.68	
18:08:06.843	training step 38920	Loss: 5.95 Examples/sec: 449.10 Hit@1: 0.88 PERR: 0.74 GAP: 0.79	

Training accuracy (Hit @1) = 88%

Validation:

Cloud ML Job, yt8m_eval_20171017_150058	All logs	Any log level	Jump to date
2017-10-17 EDT			
View Options			
15:38:23.062	examples_processed: 93824	global_step 45010 Batch Hit@1: 0.867 Batch PERR: 0.714 Batch Loss: 6.803 Examples_per_sec: 64.409	
15:38:25.363	examples_processed: 93952	global_step 45010 Batch Hit@1: 0.805 Batch PERR: 0.645 Batch Loss: 6.996 Examples_per_sec: 65.434	

Validation accuracy (Hit @1) = 80.5%

4. LSTM with only video features (4 hr 46 min)

Training :

Cloud ML Job, yt8m_train_20171015_185430	All logs	Any log level	Jump to date
2017-10-15 EDT			
View Options			
23:14:11.197	training step 4749	Loss: 30.71 Examples/sec: 40.21	
23:14:14.437	training step 4750	Loss: 35.09 Examples/sec: 40.30 Hit@1: 0.25 PERR: 0.14 GAP: 0.07	

Training accuracy(Hit @1) = 25%

Validation:

Cloud ML Job, yt8m_eval_20171017_155210	All logs	Any log level	Jump to date
2017-10-17 EDT			
View Options			
16:13:28.480	examples_processed: 55680	global_step 5213	Batch Hit@1: 0.266
16:13:30.798	examples_processed: 55808	global_step 5213	Batch Hit@1: 0.289

Validation accuracy(Hit @1) = 28.9%

5. LSTM with video and audio features (4 hr 50 min)

Training:

Cloud ML Job, yt8m_train_20171015_185004	All logs	Any log level	Jump to date
2017-10-15 EDT			
View Options			
23:15:28.017	training step 4749	Loss: 30.27 Examples/sec: 39.35	
23:15:31.388	training step 4750	Loss: 30.53 Examples/sec: 39.49 Hit@1: 0.41 PERR: 0.22 GAP: 0.17	

Training accuracy(Hit @1) = 41%

Validation:

Cloud ML Job, yt8m_eval_20171017_154855	All logs	Any log level	Jump to date
2017-10-17 EDT			
View Options			
16:11:24.939	examples_processed: 68736	global_step 5102	Batch Hit@1: 0.328
16:11:27.421	examples_processed: 68864	global_step 5102	Batch Hit@1: 0.391

Validation accuracy(Hit @1) = 39.1%

6. GRU with only video features (train duration:1 day 16hrs)

Training:

Cloud ML Job, yt8m_train_20171015_034213	All logs	Any log level	Jump to date
2017-10-16 EDT			
View Options			
20:04:20.925	training step 49409	Loss: 3134.08 Examples/sec: 45.43	
20:04:23.819	training step 49410	Loss: 3132.55 Examples/sec: 45.54 Hit@1: 0.26 PERR: 0.06 GAP: 0.00	

Training accuracy (Hit @1) = 26%

Validation:

Cloud ML Job, yt8m_eval_20171017_160421 ▾ All logs ▾ Any log level ▾ Jump to date ▾

2017-10-17 EDT							View Options ▾
▶	i	16:23:48.374	examples_processed: 56448		global_step 49575		Batch Hit@1: 0.211 Batch PERR: 0.058 Batch Loss: 3091.086... ⋮
▶	i	16:23:50.790	examples_processed: 56576		global_step 49575		Batch Hit@1: 0.273 Batch PERR: 0.061 Batch Loss: 3077.235... ⋮

Validation accuracy (Hit @1) = 27..3%

7. GRU with video and audio features (train duration:18 hr 51 min)

Training:

Cloud ML Job, yt8m_train_20171016_020609 ▾ All logs ▾ Any log level ▾ Jump to date ▾

2017-10-16 EDT							View Options ▾
▶	i	20:58:44.428	training step 22769		Loss: 2034.24	Examples/sec: 44.95	⋮
▶	i	20:58:47.370	training step 22770		Loss: 2040.13	Examples/sec: 44.63 Hit@1: 0.10 PERR: 0.00 GAP: 0.00	⋮

Training accuracy(Hit @1) = 10%

Validation:

Cloud ML Job, yt8m_eval_20171017_160652 ▾ All logs ▾ Any log level ▾ Jump to date ▾

2017-10-17 EDT							View Options ▾
▶	i	16:26:45.042	examples_processed: 51840		global_step 22710		Batch Hit@1: 0.117 Batch PERR: 0.000 Batch Loss: 1994.120... ⋮
▶	i	16:26:47.467	examples processed: 51968		global step 22710		Batch Hit@1: 0.148 Batch PERR: 0.001 Batch Loss: 1994.390... ⋮

Validation accuracy (Hit @1) = 14.8%

8. BiLSTM with audio and video features (train duration :6 hr 15 min)

Training:

Cloud ML Job, yt8m_train_20171017_054008 ▾ All logs ▾ Any log level ▾ Jump to date ▾

2017-10-17 EDT							View Options ▾
▶	i	11:56:16.864	training step 6209		Loss: 47.17	Examples/sec: 36.61	⋮
▶	i	11:56:20.405	training step 6210		Loss: 51.13	Examples/sec: 36.65 Hit@1: 0.17 PERR: 0.04 GAP: 0.01	⋮

Training accuracy (Hit @1) = 17%

Validation:

Cloud ML Job, yt8m_eval_20171017_155428 ▾ All logs ▾ Any log level ▾ Jump to date ▾

2017-10-17 EDT							View Options ▾
▶	i	16:17:37.322	examples_processed: 53504		global_step 6195		Batch Hit@1: 0.172 Batch PERR: 0.048 Batch Loss: 45.555 Examples_per_sec: 58.086 ⋮
▶	i	16:17:39.725	examples_processed: 53632		global_step 6195		Batch Hit@1: 0.180 Batch PERR: 0.050 Batch Loss: 47.444 Examples_per_sec: 60.610 ⋮

Validation accuracy (Hit @1) = 18%

9. BiGRU with audio and video features (train duration:8 hr 59 min)

Training:

Cloud ML Job, yt8m_train_20171017_025719	All logs	Any log level	Jump to date
2017-10-17 EDT			
View Options			
11:56:43.200	training step 10159	Loss: 966.41 Examples/sec: 39.08	
11:56:46.459	training step 10160	Loss: 966.41 Examples/sec: 39.96 Hit@1: 0.20 PERR: 0.05 GAP: 0.00	

Training accuracy(Hit @1) = 20%

Validation:

Cloud ML Job, yt8m_eval_20171017_155515	All logs	Any log level	Jump to date
2017-10-17 EDT			
View Options			
16:20:22.547	examples_processed: 60032	global_step 10010	Batch Hit@1: 0.109 Batch PERR: 0.027 Batch Loss: 902.370 Examples_per_sec: 54.338
16:20:25.064	examples_processed: 60160	global_step 10010	Batch Hit@1: 0.250 Batch PERR: 0.076 Batch Loss: 896.163 Examples_per_sec: 58.816

Validation accuracy (Hit @1) = 25%

5.1 Comparison

From the facts observed from the various research papers for this problem, better accuracies can be obtained if

1. Training is done longer i.e for 3 to 5 days using multiple clusters
2. Ensemble of five or seven different models are used

6 Problems and Possible improvements:

1. GRU with two layers and layer size = 2048 would give better results but when this method was attempted , it resulted in resource exhaustion error on the GCP
2. BI-LSTM when coupled with batch normalization could possibly improve the accuracy from the present value
3. During the implementation of BiLSTM if 'state_is_tuple' is not set with false , it results in tensor shape mismatch error . This fact and the detail for concatenating the forward and backward states was referred from the discussion thread of YouTube-8m competition in Kaggle
4. Stacked BiLSTM tensorflow could not be used as the 3D tensor could not be split in list of 2D tensors .This list has to be supplied as input to this API .Instead of this method dynamic bidirectional rnn was used

7 Conclusion

The consideration of this problem statement helped to learn about newer paradigms in deep learning and tensorflow APIs apart from latest deep learning network architectures

8 References

- [1] <https://research.google.com/pubs/pub45611.html>
- [2] <https://people.cs.pitt.edu/~milos/courses/cs2750-Spring04/lectures/class22.pdf>
- [3] <https://www.kaggle.com/c/youtube8m/discussion>
- [4] <https://research.google.com/youtube8m/workshop.html>
- [5] <https://static.googleusercontent.com/media/research.google.com/en//youtube8m/workshop2017/c11.pdf>
- [6] <https://static.googleusercontent.com/media/research.google.com/en//youtube8m/workshop2017/c04.pdf>
- [7] <https://static.googleusercontent.com/media/research.google.com/en//youtube8m/workshop2017/c14.pdf>
- [8] <https://static.googleusercontent.com/media/research.google.com/en//youtube8m/workshop2017/c03.pdf>